

# Hyperparameter Tuning, Regularization, Optimization

Sidharth Baskaran

July 2021

## Tuning Process

- Choose parameter combinations (points) at random in a space
- 2 parameters - grid
- 3 parameters - cube
- Coarse to fine sampling
  - Sample densely within a region with better parameter combinations
  - Coarse sample of all combinations

## Scale to pick hyperparameters

- Use a log scale to range  $\alpha$  for example
  - Set  $r \in [a, b]$
  - Then  $\alpha \in [10^a, 10^b]$  satisfying the range
- Hyperparameters for exponentially weighted averages
  - Set  $r \in [a, b]$
  - $1 - \beta = 10^r$
  - $\beta = 1 - 10^r$
  - Sampling regime important, as small changes change the values in Adam optimization largely

## Pandas vs. Caviar

- Babysitting one model
  - Adjust hyperparameters one day at a time while observing performance
  - Can't train many models at once
  - Panda approach
- Train many models in parameter with different hyperparameter settings
  - Caviar approach

## Normalizing Activations

- $\mu = \frac{1}{M} \sum_i x^{(i)}$
- $x = x - \mu$
- $\sigma^2 = \frac{1}{m} \sum_i x^{2(i)}$
- $x = x / \sigma$
- Makes contours more circular, beneficial for gradient descent
- Batch normalization  $\rightarrow$  normalize  $a^{[2]}$  to train  $w^{[3]}, b^{[3]}$  faster
- Normalize the values of  $z^{[2]}$  instead of  $a^{[2]}$
- Implementing Batch Norm
  - Given intermediate NN values  $z^{(1)}, \dots, z^{(i)}$




Figure 1: Batch norm

- \*  $\mu = \frac{1}{M} \sum_i x^{(i)}$
- \*  $\sigma^2 = \frac{1}{m} \sum_i (z_i - \mu)^2$
- \*  $z_{\text{norm}}^{(i)} = \frac{z^{(i)} - \mu}{\sqrt{\sigma^2 + \epsilon}}$
- $\bar{z}^{(i)} = \gamma z_{\text{norm}}^{(i)} + \beta$ 
  - \* These are hyperparameters which allow for adjustment of  $z_{\text{norm}}^{(i)}$
- Use  $\bar{z}^{[l](i)}$  instead of  $z^{[l](i)}$

## Fitting Batch Norm into network

- Example computation for first layer

$$x \xrightarrow{w^{[1]}, b^{[1]}} z^{[1]} \xrightarrow{\beta^{[1]}, \gamma^{[1]}} \bar{z}^{[1]} \rightarrow a^{[1]} = g^{[1]}(\bar{z}^{[1]})$$

- Repeat for each layer
- Parameters are  $W^{[k]}, b^{[k]}, \beta^{[k]}, \gamma^{[k]}$  for  $k \in [1, L]$  are the parameters
- Working with mini-batches
  - Conduct batch norm for each minibatch (i.e. starting with  $X^{\{1\}}$ )
  - Batch norm only uses examples from the current minibatch
- $b^{[l]}$  gets subtracted out in the mean subtraction step, so **eliminate this parameter, is redundant**
  - Simply remove from the subtraction
- Implementing gradient descent
  - For  $t \in [1, \text{numMiniBatch}]$ 
    - \* Compute forward prop on  $X^{\{t\}}$ 
      - In each hidden layer, use batch norm
    - \* Use backprop to calculate  $dw^{[l]}, d\beta^{[l]}, d\gamma^{[l]}$  and update parameters

## Batch norm reasoning

- Conducts input normalization but for hidden units
- Makes deeper weights robust to changes in weights of earlier layers
- Learning on shifting input distribution  $\rightarrow$  covariant shift

- If  $X$  distribution changes, must retrain for  $X \rightarrow Y$  mapping, so values input to a hidden layer change
- Batch norm reduces this effect, layers are more independent
- Batch norm as regularization
  - Each minibatch scaled by mean/variance of the examples in it
  - Adds noise to  $z^{[l]}$ , so adds noise to hidden layer activations similar to dropout
  - Has slight regularization effect, but with a larger minibatch size this is reduced

## Batch Norm at Test Time

- Need to estimate  $\mu, \sigma^2$  at test time using weighted average across minibatches
- Calculate  $\mu^{\{1\}^l}, \dots, \mu^{\{n\}^l}$ , using a moving average, same for  $\sigma^2$