

Hyperparameter Tuning, Regularization, Optimization

Sidharth Baskaran

July 2021

Tuning Process

- Choose parameter combinations (points) at random in a space
- 2 parameters - grid
- 3 parameters - cube
- Coarse to fine sampling
 - Sample densely within a region with better parameter combinations
 - Coarse sample of all combinations

Scale to pick hyperparameters

- Use a log scale to range α for example
 - Set $r \in [a, b]$
 - Then $\alpha \in [10^a, 10^b]$ satisfying the range
- Hyperparameters for exponentially weighted averages
 - Set $r \in [a, b]$
 - $1 - \beta = 10^r$
 - $\beta = 1 - 10^r$
 - Sampling regime important, as small changes change the values in Adam optimization largely

Pandas vs. Caviar

- Babysitting one model
 - Adjust hyperparameters one day at a time while observing performance
 - Can't train many models at once
 - Panda approach
- Train many models in parameter with different hyperparameter settings
 - Caviar approach

Normalizing Activations

- $\mu = \frac{1}{M} \sum_i x^{(i)}$
- $x = x - \mu$
- $\sigma^2 = \frac{1}{m} \sum_i x^{2(i)}$
- $x = x / \sigma$
- Makes contours more circular, beneficial for gradient descent
- Batch normalization \rightarrow normalize $a^{[2]}$ to train $w^{[3]}, b^{[3]}$ faster
- Normalize the values of $z^{[2]}$ instead of $a^{[2]}$
- Implementing Batch Norm
 - Given intermediate NN values $z^{(1)}, \dots, z^{(i)}$




Figure 1: Batch norm

- * $\mu = \frac{1}{M} \sum_i x^{(i)}$
- * $\sigma^2 = \frac{1}{m} \sum_i (z_i - \mu)^2$
- * $z_{\text{norm}}^{(i)} = \frac{z^{(i)} - \mu}{\sqrt{\sigma^2 + \epsilon}}$
- $\bar{z}^{(i)} = \gamma z_{\text{norm}}^{(i)} + \beta$
 - * These are hyperparameters which allow for adjustment of $z_{\text{norm}}^{(i)}$
- Use $\bar{z}^{[l](i)}$ instead of $z^{[l](i)}$

Fitting Batch Norm into network

- Example computation for first layer

$$x \xrightarrow{w^{[1]}, b^{[1]}} z^{[1]} \xrightarrow{\beta^{[1]}, \gamma^{[1]}} \bar{z}^{[1]} \rightarrow a^{[1]} = g^{[1]}(\bar{z}^{[1]})$$

- Repeat for each layer
- Parameters are $W^{[k]}, b^{[k]}, \beta^{[k]}, \gamma^{[k]}$ for $k \in [1, L]$ are the parameters
- Working with mini-batches
 - Conduct batch norm for each minibatch (i.e. starting with $X^{\{1\}}$)
 - Batch norm only uses examples from the current minibatch
- $b^{[l]}$ gets subtracted out in the mean subtraction step, so **eliminate this parameter, is redundant**
 - Simply remove from the subtraction
- Implementing gradient descent
 - For $t \in [1, \text{numMiniBatch}]$
 - * Compute forward prop on $X^{\{t\}}$
 - In each hidden layer, use batch norm
 - * Use backprop to calculate $dw^{[l]}, d\beta^{[l]}, d\gamma^{[l]}$ and update parameters

Batch norm reasoning

- Conducts input normalization but for hidden units
- Makes deeper weights robust to changes in weights of earlier layers
- Learning on shifting input distribution \rightarrow covariant shift

- If X distribution changes, must retrain for $X \rightarrow Y$ mapping, so values input to a hidden layer change
- Batch norm reduces this effect, layers are more independent
- Batch norm as regularization
 - Each minibatch scaled by mean/variance of the examples in it
 - Adds noise to $z^{[l]}$, so adds noise to hidden layer activations similar to dropout
 - Has slight regularization effect, but with a larger minibatch size this is reduced

Batch Norm at Test Time

- Need to estimate μ, σ^2 at test time using weighted average across minibatches
- Calculate $\mu^{\{1\}l}, \dots, \mu^{\{n\}l}$, using a moving average, same for σ^2

Softmax regression

- If there are multiple classes (e.g. Chick, Dog, Cat with classes 3, 2, 1, and 0 if none)
 - Indexes are $i \in [0, C - 1]$ with C classes




Figure 2: Softmax NN

- Output is a 4-dim layer, each node contains probability $P(\text{class}|x)$
 - Sum must be 1
- Softmax output layer generates these outputs
 - Is an activation function
- Compute a temp. var $t = e^{z^{[L]}}$ elementwise of dim (4,1)
 - $a^{[L]} = \frac{e^{z^{[L]}}}{\sum_{i=1}^4 t_i}$ such that $a_i^{[L]} = \frac{t_i}{\sum_{i=1}^4 t_i}$
- Resulting probabilities are the desired classifications
- Unusual as it takes in vectorial input, not some \mathbb{R}^f
- Softmax can represent linear decision boundaries between multiple classes (without hidden layers to learn nonlinear ones)

Training a Softmax Classifier

- A hardmax would map the vector $z^{[L]}$ to either 0 or 1, whereas softmax is a gentle, more precise mapping
- Is a generalization of logistic regression to C classes
- Loss function
 - Makes corresponding probability of the ground truth class as high as possible \rightarrow minimizes the loss

$$a^{[u]} = g^{[L]}(z^{[L]}) = \begin{bmatrix} e^5/(e^5 + e^2 + e^{-1} + e^3) \\ e^2/(e^5 + e^2 + e^{-1} + e^3) \\ e^{-1}/(e^5 + e^2 + e^{-1} + e^3) \\ e^3/(e^5 + e^2 + e^{-1} + e^3) \end{bmatrix} = \begin{bmatrix} 0.842 \\ 0.042 \\ 0.002 \\ 0.114 \end{bmatrix}$$

Figure 3: Example

$$\mathcal{L}(\hat{y}, y) = - \sum_{j=1}^4 y_j \log \hat{y}_j$$

- Cost J on entire training set

$$J = \frac{1}{m} \sum_{i=1}^m \mathcal{L}(\hat{y}^{(i)}, y^{(i)})$$

- Gradient descent with softmax
 - Key equation is $dz^{[L]} = \hat{y} - y$ which are all (4,1) vectors

Deep learning frameworks

- Ease of programming
- Running speed
- Open source

TensorFlow

- Motivating problem
 - Want to minimize $J(w) = w^2 - 10w + 25 = (w - 5)^2$

```
import numpy as np
import tensorflow as tf

# initialize to 0 float32
w = tf.Variable(0, dtype=tf.float32)
# alpha=0.1
optimizer = tf.keras.optimizers.Adam(0.1)

# one training step iteration
def train_step():
    with tf.GradientTape() as tape:
        cost = w**2 - 10*w + 25
    trainable_variables = w
    grads = tape.gradient(cost, trainable_variables)
    optimizer.apply_gradients(zip(grads, trainable_variables))
```

- Gradient tape records order of operations, thus only need to implement forward prop step
- Can also implement this using optimization

```

w = tf.Variable(0, dtype=tf.float32)
x = np.array([1.0,-10.0,25.0], dtype=np.float32)
optimizer = tf.keras.optimizers.Adam(0.1)

def training(x,w,optimizer):
    def cost_fn():
        return x[0]*w**2 + x[1]*w + x[2]
    for i in range(1000):
        optimizer.minimize(cost_fn, [w])
    return w

w = training(x,w,optimizer)

```