

ML Strategy 1

Sidharth Baskaran

July 2021

Orthogonalization

- TV tuning example
 - Multiple knobs to transform image in a specific way
 - Orthogonalization - each knob has one function without affecting others
- Easier to tune
- Chain of assumptions in ML
 - Fit training set well on cost function
 - Fit dev set well on cost function
 - Fit test set well on cost function
 - Performs well in real world
- On each assumption, want a distinct set of knobs to tune # Single Number Evaluation Metric
- Precision \rightarrow of classified examples, which were classified correctly
- Recall \rightarrow of all images of some class, how many were correctly classified
- Define anew metric that combines P and R
 - F1 score is the harmonic mean $\frac{2}{1/P+1/R}$
- Well defined dev set + single number evaluation metric allows for deciding which classifier is better

Satisficing and optimizing metrics

- Example: accuracy and running time
 - Let cost = accuracy - 0.5*running_time
- If we want to maximize accuracy but have running time at a bare minimum, then we **optimize** accuracy and **satisfice** running time
- If there are N metrics, want to optimize 1 and satisfice $N - 1$

Train/Dev/Test Set Distributions

- Dev and test sets should come from same distribution
- Dev set + metric = a target on which to aim (find best classifier)
 - Iterate to the center
- Should not generalize to one distribution over another
 - Choose dev + test set to reflect data expected in feature, important to generalize on
- Randomly shuffle to produce dev + test

Size of Dev and Test Sets and When to Change

- Size of test set → big enough to give high confidence in overall model performance
- Changing metric = changing position of target board
- A metric can indicate well, but algorithm could not perform as intended → change metric
- Error in dev set
 - $\text{Err} = \frac{1}{m_{\text{dev}}} \sum_{i=1}^{m_{\text{dev}}} \mathcal{I}\{y_{\text{pred}}^{(i)} \neq y^{(i)}\}$
 - Counts misclassified images
 - Instead, could give greater weight to images that aren't desired but prone to being classified
- Orthogonalization
 - Define metric
 - Worry separately how to do well on metric
- **If doing well on metric + dev set \neq doing well on application, change metric/dev set**

Human-level performance

- Algorithm stabilizes in accuracy after exceeding human performance → theoretical Bayes optimal error
- As long as ML worse than humans
 - Get labeled data from humans
 - Gain insight from manual error analysis
 - Better analysis of bias/variance \neq Avoidable bias
- If training and dev error much more than human error, then is a high bias problem
- If human error is similar, focus on variance
- Treat human level error as proxy for Bayes error
- Difference between Bayes (Human) error and training error is **avoidable bias**
- Difference between training and dev error is **variance**

Understanding human-level performance

- Human error is a proxy for Bayes optimal error
- If training and dev error is very similar, then want a lower human error
 - Allows for adjusting addressing avoidable bias problem which would be otherwise ignored

Surpassing Human-Level Performance

- Given multiple error sources, sometimes can't tell if it is a bias or variance issue
- Cannot rely on intuition to determine what changes to make if model surpasses human error at its lowest
- Surpassing examples
 - Online advertising
 - Product recommendations
 - Logistics
 - Loan approval
- Humans good at natural perception tasks in general

Improving Model Performance

- Fundamental assumptions of supervised learning
 - Can fit training set well
 - Training set performance generalizes well to dev/test set
- Reducing avoidable bias
 - Bigger model

Human-level error

(proxy for Bayes error)

Training error

Dev error



Figure 1: Human-level performance bias/variance

- Better/longer optimization algorithms
- NN architecture/hyperparameter search improvements
- Variance improvements
 - More data
 - Regularization $\rightarrow L_2$, dropout, data aug.
 - Better architecture