

Foundations of Convolutional Neural Networks

Sidharth Baskaran

September 2021

Computer Vision

- Image classification
 - Binary classification
- Object detection
 - Drawing boxes/bounding the objects
 - Multiple instances of object
- Neural Style Transfer
 - Content and style images
 - Repaint content w/ style
- Inputs can get large
 - E.g 64x64x3 is small but larger images have many input features




Figure 1: Example network

- If $x \in \mathbb{R}^{3M}$ then $W^{[1]} \rightarrow (1000, 3M)$
 - $W^{[1]}x$ gives the output network vector of dimension $(1000, 1)$
- Implementing the convolution operation is efficient for large input images

Edge Detection

- Detecting certain feature sets of images
 - E.g. vertical and horizontal edges

- Given a 6x6 grayscale image, apply a 3x3 kernel or filter
 - Convolution operator** * convolves filter over image
- Paste filter on the first such region of the image, and take elementwise product and then sum to obtain value
- Output a 4x4 image




Figure 2: Basic Convolution

- Shift the kernel stepwise to the left to fill up the output
- Is 4x4 as can shift downwards/left/right to obtain 4 unique locations
 - $\dim(\text{Im}, 1) - \dim(\text{Ker}, 1) + 1$




Figure 3: Example

- In example, detects a light to dark transition
- Can also make distinction
- Example:* Sobel filter \rightarrow puts weight towards central pixel, more robust for edge detection

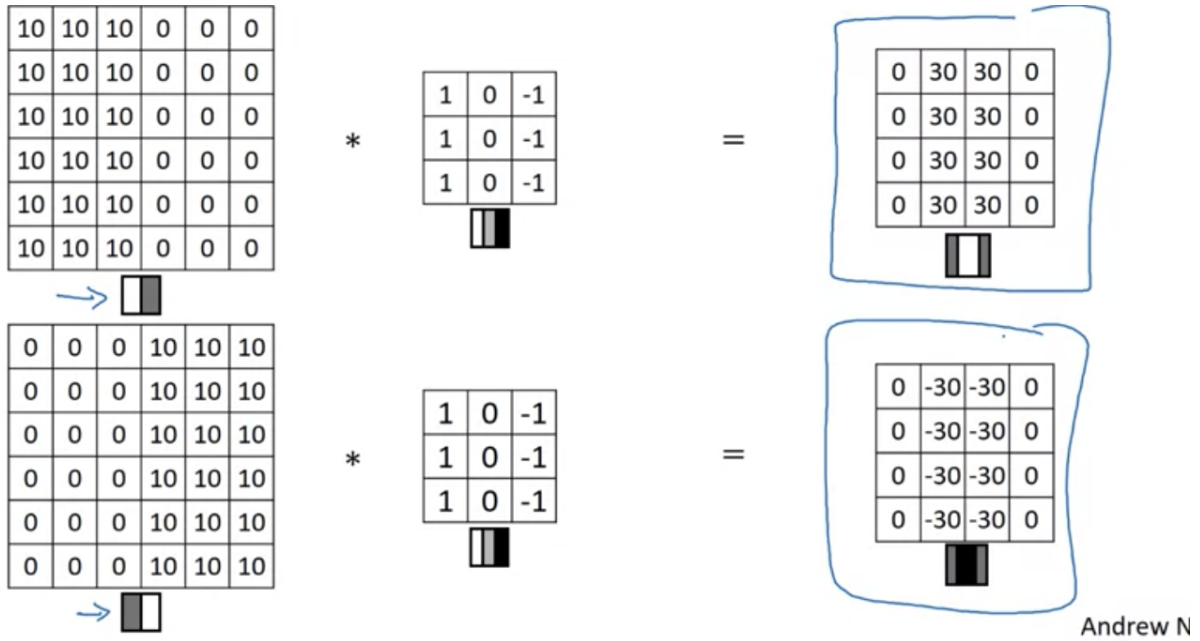


Figure 4: Transition examples

- Rotate 90 degrees for horizontal/vertical differentiation
- Can learn the filter values with backprop
 - Treat as parameters, learn

Padding

- Given an image with input dimensions $n \times n$ and filter with $f \times f$
 - Output dimensions are $n - f + 1 \times n - f + 1$
- Downsides \rightarrow image shrinks
 - Lots of overlap on central pixels but corner pixels represented less
- Can pad image with border of 1px, for example
 - $6 \times 6 \rightarrow 8 \times 8$ image
- Can preserve dimension of output image
 - Padding pixels are 0 by convention
 - Let $p = 1$ be padding amount
- Valid and Same convolutions
 - Valid \rightarrow no padding
 - $n \times n * f \times f \rightarrow n - f + 1 \times n - f + 1$
 - Same \rightarrow pad s.t. output size = input size
- By convention, f is always odd, i.e $f \bmod 2 \neq 0$